


Gap-filling loss curves at best config ($h = 64$, $\text{lr} = 0.05$, $n_{\text{train}} = 200$, 3000 epochs, 10 seeds, $\pm 1\sigma$)

Training loss (MSE)



Validation loss (MSE)

